

**PHCM9795**

**Foundations of Biostatistics**

**Term 2, 2026**





# Table of contents

|                                                               |           |
|---------------------------------------------------------------|-----------|
| <b>Course introduction</b>                                    | <b>1</b>  |
| Welcome                                                       | 1         |
| Course information                                            | 1         |
| Course aim                                                    | 1         |
| Learning outcomes                                             | 1         |
| Suggested reference                                           | 2         |
| Change log                                                    | 2         |
| <b>1 Introduction to biostatistics and data fundamentals</b>  | <b>3</b>  |
| Learning objectives                                           | 3         |
| Optional readings                                             | 3         |
| 1.1 An introduction to statistics                             | 3         |
| 1.2 What are data?                                            | 4         |
| 1.3 Descriptive and inferential statistics                    | 5         |
| 1.4 Summarising continuous data                               | 6         |
| 1.5 Exploratory data analysis for continuous data             | 11        |
| <b>An introduction to jamovi</b>                              | <b>13</b> |
| 1.6 An introduction to jamovi                                 | 13        |
| Learning outcomes                                             | 13        |
| 1.7 Introduction                                              | 13        |
| 1.8 Part 1: An introduction to jamovi                         | 13        |
| 1.9 Installing jamovi                                         | 13        |
| 1.10 A simple jamovi analysis                                 | 14        |
| 1.11 Part 2: Obtaining summary statistics for continuous data | 20        |
| 1.12 Setting a value to missing                               | 23        |
| <b>An introduction to R and RStudio</b>                       | <b>27</b> |
| Learning outcomes                                             | 27        |
| 1.13 Part 1: An introduction to R                             | 27        |
| 1.14 Part 2: Obtaining summary statistics for continuous data | 41        |
| 1.15 Setting a value to missing                               | 50        |
| <b>Activities</b>                                             | <b>53</b> |
| <b>Appendix</b>                                               | <b>55</b> |
| <b>Analysis flowchart</b>                                     | <b>57</b> |
| <b>References</b>                                             | <b>59</b> |



# Course introduction

## Welcome

This introductory course in biostatistics aims to provide students with core biostatistical skills to analyse and present quantitative data from different study types. These are essential skills required in your degree and throughout your career.

We hope you enjoy the course and will value your feedback and comment throughout the course.

## Course information

Biostatistics is a foundational discipline needed for the analysis and interpretation of quantitative information and its application to population health policy and practice.

This course is central to becoming a population health practitioner as the concepts and techniques developed in the course are fundamental to your studies and practice in population health. In this course you will develop an understanding of, and skills in, the core concepts of biostatistics that are necessary for analysis and interpretation of population health data and health literature.

In designing this course, we provide a learning sequence that will allow you to obtain the required graduate capabilities identified for your program. This course is taught with an emphasis on formulating a hypothesis and quantifying the evidence in relation to a specific research question. You will have the opportunity to analyse data from different study types commonly seen in population health research.

The course will allow those of you who have covered some of this material in your undergraduate and other professional education to consolidate your knowledge and skills. Students exposed to biostatistics for the first time may find the course challenging at times. Based on student feedback, the key to success in this course is to devote time to it every week. We recommend that you spend an average of 10-15 hours per week on the course, including the time spent reading the course notes and readings, listening to lectures, and working through learning activities and completing your assessments. Please use the resources provided to assist you, including online support.

## Course aim

This course aims to provide students with the core biostatistical skills to apply appropriate statistical techniques to analyse and present population health data.

## Learning outcomes

On successful completion of this course, you will be able to:

1. Summarise and visualise data using statistical software.
2. Demonstrate an understanding of statistical inference by interpreting p-values and confidence intervals.
3. Apply appropriate statistical tests for different types of variables given a research question, and interpret computer output of these tests appropriately.
4. Determine the appropriate sample size when planning a research study.
5. Present and interpret statistical findings appropriate for a population health audience.

### Suggested reference

UNSW Sydney, School of Population Health. PHCM9795: Foundations of Biostatistics Course Notes. UNSW Sydney, 2026.

### Change log

1 August

- Section 10.12. Corrected typos: Based on the information in Worked Example ~~10.5~~ 10.3, the Power is 0.9 and Significance level is 0.05, and the two proportions are entered as 0.35 and ~~0.2~~ 0.25. Define equal numbers in each group (Ratio of 1).

27 July

- corrected the t-test in Section 9.2 to use a Welch's t-test

14 July

- fixed the ordering of the categories in Table 7.5

9 June 2026

- updated the address of the Binomial Distribution applet (Module 2) and the Normal Distribution applet (Module 3)
- clarified the instructions for Activity 2.5

1 June 2026

- Section 1.15 - corrected typo: "We could use *similar* syntax..." Also, the code `sample$weight_clean = ifelse(sample$weight==700.2, NA, sample$weight)` should be correctly written as `sample$weight_clean = ifelse(sample$weight==700.2, 70.2, sample$weight)`.

# Module 1

## Introduction to biostatistics and data fundamentals

### Learning objectives

By the end of this module, you will be able to:

- Understand the difference between descriptive and inferential statistics
- Distinguish between different types of variables
- Present and report continuous data numerically and graphically
- Compute summary statistics to describe the centre and spread of data

### Optional readings

Kirkwood and Sterne (2001); Chapters 2, 3 and 4. [\[UNSW Library Link\]](#)

Bland (2015); Chapter 4. [\[UNSW Library Link\]](#)

### 1.1 An introduction to statistics

The dictionary of statistics (Upton and Cook, 2008) defines statistics simply as: “The science of collecting, displaying, and analysing data.”

Statistics is a branch of mathematics, and there are two main divisions within the field of statistics: mathematical statistics and applied statistics. Mathematical statistics deals with development of new methods of statistical inference and requires detailed knowledge of abstract mathematics for its implementation. Applied statistics applies the methods of mathematical statistics to specific subject areas, such as business, psychology, medicine and sociology.

Biostatistics can be considered as the “application of statistical techniques to the medical and health fields”. However, biostatistics sometimes overlaps with mathematical statistics. For instance, given a certain biostatistical problem, if the standard methods do not apply then existing methods must be modified to develop a new method.

### Scope of Biostatistics

Research is essential in the practice of health care. Biostatistical knowledge helps health professionals in deciding whether to prescribe a new drug for the treatment of a disease or to advise a patient to give up drinking alcohol. To practice evidence-based healthcare, health professionals must keep abreast of the latest research, which requires understanding how the studies were designed, how data were collected and analysed, and how the results were interpreted. In clinical medicine, biostatistical methods are used to determine the accuracy of a measurement, the efficacy of a drug in treating a disease, in comparing different measurement techniques, assessing diagnostic tests, determining normal values, estimating prognosis and

monitoring patients. Public health professionals are concerned about the administration of medical services or ensuring that an intervention program reduces exposure to certain risk factors for disease such as life-style factors (e.g. smoking, obesity) or environmental contaminants. Knowledge of biostatistics helps determine them make decisions by understanding, from research findings, whether the prevalence of a disease is increasing or whether there is a causal association between an environmental factor and a disease.

The value of biostatistics is to transform (sometimes vast amounts of) data into meaningful information, that can be used to solve problems, and then be translated into practice (i.e. to inform public health policy and decision making). When undertaking research having a biostatistician as part of a multidisciplinary team from the outset, together with scientists, clinicians, epidemiologists, healthcare specialists is vital, to ensure the validity of the research being undertaken and that information is interpreted appropriately.

## 1.2 What are data?

According to the Australian Bureau of Statistics, “data are measurements or observations that are collected as a source of information”.<sup>1</sup> Note that technically, the word *data* is a plural noun. This may sound a little odd, but it means that we say “data are ...” when discussing a set of measurements.

Other definitions that we use in this course are:

- **observation**, (or **record**, or **unit record**): one individual in the population being studied
- **variable**: a characteristic of an individual being measured. For example, height, weight, eye colour, income, country of birth are all types of variables.
- **dataset**: the complete collection of all observations

### Types of variables

We can categorise variables into two main types: numeric or categorical.

**Numerical variables** (also called quantitative variables) comprise data that must be represented by a number, which can be either measured or counted.

**Continuous** variables can take any value within a defined range.

For example, age, height, weight or blood pressure, are continuous variables because we can make any divisions we want on them, and they can be measured as small as the instrument allows. As an illustration, if two people have the same blood pressure measured to the nearest millimetre of mercury, we may get a difference between them if the blood pressure is measured to the nearest tenth of millimetre. If they are still the same (to the nearest tenth of a millimetre), we can measure them with even finer gradations until we can see a difference.

**Discrete** variables can only take one of a distinct set of values (usually whole numbers). For discrete variables, observations are based on a quantity where both ordering and magnitude are important, such that numbers represent actual measurable quantities rather than mere labels.

For example, the number of cancer cases in a specified area emerging over a certain period, the number of motorbike accidents in Sydney, the number of times a woman has given birth, the number of beds in a hospital are all discrete variables. Notice that a natural ordering exists among the data points, that is, a hospital with 100 beds has more beds than a hospital with 75 beds. Moreover, a difference between 40 and 50 beds is the same as the difference between 80 and 90 beds.

**Categorical variables** comprise data that describe a ‘quality’ or ‘characteristic’. Categorical variables, sometimes called qualitative variables, do not have measurable numeric values. Categorical variables can be nominal or ordinal.

A **nominal** variable consists of unordered categories. For example, gender, race, ethnic group, religion, eye colour etc. Both the order and magnitude of a nominal variable are unimportant.

<sup>1</sup><https://www.abs.gov.au/statistics/understanding-statistics/statistical-terms-and-concepts/data>



If a nominal variable takes on one of two distinct categories, such as black or white then it is called a **binary** or dichotomous variable. Other examples would be smoker or non-smoker; exposed to arsenic or not exposed.

A nominal variable can also have more than two categories, such as blood group, with categories of: Group A, Group B, Group AB and Group O.

**Ordinal** variables consist of ordered categories where differences between categories are important, such as socioeconomic status (low, medium, high) or student evaluation rating could be classified according to their level of satisfaction: (highly satisfied, satisfied and unsatisfied). Here a natural order exists among the categories.

Note that categorical variables are often stored in data sets using numbers to represent categories. However, this is for convenience only, and these variable must not be analysed as if they were numeric variables.

### 1.3 Descriptive and inferential statistics

When analysing a set of data, it is important to consider the aims of the analysis and whether these are *descriptive* or *inferential*. Essentially, descriptive statistics summarise data from a single sample or population, and present a “snap-shot” of those data. Inferential statistics use sample data to make statements about larger populations.

#### Descriptive statistics

Descriptive statistics provide a ‘picture’ of the characteristics of a population, such as the average age, or the proportion of people born in Australia. Two common examples of descriptive statistics are reports summarising a nation’s birth statistics, and death statistics.

#### Births

The Australian Institute of Health and Welfare produces comprehensive reports on the characteristics of Australia’s mothers and babies using the most recent year of data from the National Perinatal Data Collection. The National Perinatal Data Collection comprises *all registered births* in Australia.

The most recent report, published in 2024, summarises Australian births from 2022. (Australian Institute of Health and Welfare (2024)).

One headline from the report is that “More First Nations mothers are accessing antenatal care in the first trimester (up from 51% in 2013 to 71% in 2022)”. The report presents further descriptive statistics, such as the average maternal age (31.2 years) and the proportion of women giving birth by caesarean (39%).

#### Deaths

In another example, consider characteristics of all deaths in Australia in 2023 (Australian Bureau of Statistics (2024)).

“COVID-19 was the ninth leading cause of death in 2023, after ranking third in 2022.”

The report presents the leading causes of death in 2023:

“The leading cause of death was ischaemic heart disease, accounting for 9.2% of deaths. The gap between ischaemic heart disease and dementia (the second leading cause of death) has continued to narrow over time, with only 237 deaths separating the top two leading causes in 2023.”

The top five causes of death are also presented as a graph, enabling a simple comparison of the changes in rates of death between 2014 and 2023.




Figure 1.1: Leading causes of death, age-standardised death rates, 2014-2023

## Inferential statistics

Inferential statistics use data collected from a sample to make conclusions (inferences) about the whole population from which the sample was drawn. For example, the Australian Institute of Health and Welfare's **Australia's health** reports (eg Australian Institute of Health and Welfare (2025)) use a representative sample to make estimates of the health of the whole of Australia. We will revisit *inferential statistics* in later modules.

### 1.4 Summarising continuous data

In the first two Modules, we will focus on ways to summarise and present data. We will see that the choice of presentation will depend on the type of variable being summarised. In this Module, we will focus on continuous variables, and will focus on categorical data in Module 2.

#### Summarising a single continuous variable numerically

When summarising continuous data numerically, there are two things we want to know:

1. What is the average value? And,
2. How variable (or spread out) are the data?

We will use a sample of 35 ages (in whole years) to illustrate how to calculate the average value and measures of variability:

59 41 44 43 31 47 53 59 35 60 54 61 67 52 43 46 39 69 50 64 57 39 54 50 51 31 48 49 70 44 60 51 37 53 34

#### Measures of central tendency

##### Mean

The most commonly used measure of the central tendency of the data is the mean, calculated as the sum of all the values divided by the number of values:

$$\bar{x} = \frac{\sum x}{n}$$

From the age example:  $\bar{x} = 1745/35 = 49.9$ . Thus, the mean age of this sample is 49.9 years.

### Median

Another measure of central tendency is the median. The median is the middle value of the data, the value at which half of the measurements lie above it and half of the measurements lie below it.

To estimate the median, the data are ordered from the lowest to highest values, and the middle value is used. If the middle value is between two data points (if there are an even number of observations), the median is an average of the two values.

Using our example, we could rank the ages from smallest to largest, and locate the middle value (which has been bolded):

31 31 34 35 37 39 39 41 43 43 44 44 46 47 48 49 50 **50** 51 51 52 53 53 54 54 57 59 59 60 60 61  
64 67 69 70

Here, the median age is 50 years.

Note that, in practice, the median is usually calculated by software automatically, and there is no need to rank our data.

### Describing the spread of the data

In addition to measuring the centre of the data, we also need an estimate of the variability, or spread, of the data points.

#### Range

The absolute measure of the spread of the data is the range, that is the difference between the highest and lowest values in the dataset.

Range = highest data value – lowest data value

Using the age example, Range = 70 - 31 = 39 years.

The range is most usefully reported as the actual lowest and highest values e.g. Range: 31 to 70 years.

The range is not always ideal as it only describes the extreme values, without considering how the bulk of the data is distributed between them.

### Variance and standard deviation

More useful statistics to describe the spread of the data around a mean value are the variance and standard deviation. These measures of variability depend on the difference between individual observations and the mean value (deviations). If all values are equal to the mean there would be no variability at all, all deviations would be zero; conversely large deviations indicate greater variability.

One way of combining deviations in a single measure is to first square the deviations and then average the squares. Squaring is done because we are equally interested in negative deviations and positive deviations; if we averaged without squaring, negative and positive deviations would 'cancel out'. This measure is called the variance of the set of observations. It is 'the average squared deviation from the mean'. Because the variance is in 'square' units and not in the units of the measurement, a second measure is derived by taking the square root of the variance. This is the standard deviation (SD), and is the most commonly used measure of variability in practice, as it is a more intuitive interpretation since it is in the same units as the units of measurement.

The formula for the variance of a sample ( $s^2$ ) is:

$$s^2 = \frac{\sum (x - \bar{x})^2}{n - 1}$$

Note that the deviations are first squared before they are summed to remove the negative values; once summed they are divided by the sample size minus 1.

The sample standard deviation is the square root of the of the sample variance:

$$s = \sqrt{s^2}$$

For the age example, we would calculate the sample variance using statistical software. The sample standard deviation is estimated as:  $s = 10.47$  years.

Characteristics of the standard deviation:

- It is affected by every measurement
- It is in the same units as the measurements
- It can be converted to measures of precision (standard error and 95% confidence intervals) (Module 3)

### Interquartile range

The inter-quartile range (IQR) describes the range of measurements in the central 50% of values lie. This is estimated by calculating the values that cut the data at the bottom 25% and top 25%. The IQR is the preferred measure of spread when the median has been used to describe central tendency.

In the age example, the IQR is estimated as 43 to 58 years.

### Population values: mean, variance and standard deviation

The examples above show how the sample mean, range, variance and standard deviation are calculated from the sample of ages from 35 people. If we had information on the age of the *entire* population that the sample was drawn from, we could calculate all the summary statistics described above (for the sample) for the population.

The equation for calculating the population mean is the same as that of sample mean, though now we denote the population mean as  $\mu$ :

$$\mu = \frac{\sum x}{N}$$

Where  $\sum x$  represents the sum of the values in the population, and  $N$  represents the total number of measurements in the population.

To calculate the population variance ( $\sigma^2$ ) and standard deviation ( $\sigma$ ), we use a slightly modified version of the equation for  $s^2$ :

$$\sigma^2 = \frac{\sum (x - \mu)^2}{N}$$

with a population standard deviation of:  $\sigma = \sqrt{\sigma^2}$ .

In practice, we rarely have the information for the entire population to be able to calculate the population mean and standard deviation. Theoretically, however, these statistics are important for two main purposes:

1. the characteristics of the normal distribution (the most important probability distribution discussed in later modules) are defined by the population mean and standard deviation;
2. while calculating sample sizes (discussed in later modules) we need information about the population standard deviation, which is usually obtained from the existing literature.

### Summarising a single continuous variable graphically

As well as calculating measures of central tendency and spread to describe the characteristics of the data, a graphical plot can be helpful to better understand the characteristics and distribution of the measurements obtained. *Histograms*, *density plots* and *box plots* are excellent ways to display continuous data graphically.

### Frequency histograms

A frequency histogram is a plot of the number of observations that fall within defined ranges of non-overlapping intervals (called bins). Examples of frequency histograms are given in Figure 1.2.




Figure 1.2: Histograms of age and serum bilirubin from PBC data

Some features of a frequency histogram:

- The area under each rectangle is proportional to the frequency
- The rectangles are drawn without gaps between them (that is, the rectangles touch)
- The data are 'binned' into discrete intervals (usually of equal width)

A slight variation on the frequency histogram is the **density histogram**, which plots the density on the y-axis. The density is a technical term, which is similar to the relative frequency, but is scaled so that the sum of the area of the bars is equal to 1.

Both the frequency and density histograms are useful for understanding how the data is distributed across the range of values. Taller bars indicate regions where the data is more densely concentrated, while shorter bars represent areas with fewer data points.

### Density plot

A density plot can be thought of as a smoothed version of a density histogram. Like histograms, density plots show areas where there are a lot of observations and areas where there are relatively few observations. Figure 1.3 illustrates example density plots for the same data as plotted in Figure 1.2.

Like histograms, density plots allow you to see the overall shape of a distribution. They are most useful when there are only a small number of observations being plotted. When plotting small datasets, the shape of a histogram can depend on how the bins are defined. This is less of an issue if a density plot is used.




Figure 1.3: Density plots of age (left) and serum bilirubin (right) from a sample of data

### Boxplots

Another way to inspect the distribution of data is by using a box plot. In a box plot:

- the line across the box shows the median value
- the limits of the box show the 25-75% range (i.e. the inter-quartile range (IQR) where the middle 50% of the data lie)
- the bars (or whiskers) indicate the most extreme values (highest and lowest) that fall within 1.5 times the interquartile range from each end of the box
  - the upper whisker is the highest value falling within 75th percentile plus  $1.5 \times \text{IQR}$
  - the lower whisker is the lowest value falling within 25th percentile minus  $1.5 \times \text{IQR}$
- any values in the dataset lying outside the whiskers are plotted individually, often labelled by the row number of the observation.

Figure 1.4 presents two example boxplots for age and serum bilirubin.

### The shape of a distribution

Histograms and density plots allow us to consider the shape of a distribution, and in particular, whether a distribution is *symmetric* or *skewed*.

In a histogram, if the rectangles fall in a roughly symmetric shape around a single midpoint, we say that the distribution is symmetric. Similarly, if a density plot looks roughly symmetric around a single point, the distribution is symmetric.

If the histogram or density plot has a longer tail to the right, then the data are said to be positively skewed (or skewed to the right); if the histogram or density plot has an extended tail to the left, then the data are negatively skewed (or skewed to the left).

The skewness of a distribution is defined by the location of the longer tail in a histogram or density plot, not the location of the peak of the data.

From Figure 1.2 and Figure 1.3, we can see that the distribution for age is roughly symmetric, while the distribution for serum bilirubin is highly positively skewed (or skewed to the right).

While it is technically possible to determine the shape of a distribution using a boxplot, a histogram or density plot gives a more complete illustration of a distribution and would be the preferred method of assessing shape.




Figure 1.4: Box plot of age and serum bilirubin from PBC study data

### Which measure of central tendency to use

We introduced the mean and median in Section 1.4 as measures of central tendency. We need to assess the shape of a distribution to answer which is the more appropriate measure to use.

If a distribution is symmetric, the mean and median will be approximately equal. However, the mean is the preferred measure of central tendency as it makes use of every data point, and has more useful mathematical properties.

The mean is not a good measure of central tendency for skewed distributions, as the calculation will be influenced by the observations in the tail of the distribution. The median is the preferred statistic for describing central tendency in a skewed distribution.

If the data exhibits a symmetric distribution, we use the standard deviation as the measure of spread. Otherwise, the interquartile range is preferred.

## 1.5 Exploratory data analysis for continuous data

Before conducting any formal analysis, it is good practice to undertake exploratory data analysis. This analysis step gives you a high-level overview of your data: what do your data look like, what are the main features, what shape is the distribution, and are there any unusual values.

Exploring continuous data is best done by examining plots: specifically density plots and boxplots. Density plots are useful in determining the shape of a distribution, to help you decide what summary measures to use, and what type of analysis to conduct. Boxplots can be useful in identifying any unusual points - often called outliers. Outliers can be problematic and the decision to include them or omit them from further analyses can be difficult.

After detecting any outliers or extreme values, **do not automatically exclude them from the analysis**. First, it is important to check the original data collection form or questionnaire to rule out the possibility of a data entry error. If the outlier is not a data entry error, it is then important to decide whether the observation is biologically possible. This step will usually need to be answered by a topic expert. If the outlier is biologically possible, it must be included in the analysis. Only if the outlier is biologically impossible should it be set to missing.





# An introduction to jamovi

## 1.6 An introduction to jamovi

### Learning outcomes

By the end of these notes, you will be able to:

- navigate the jamovi interface
- input and import data into jamovi
- use jamovi menus to summarise data
- perform basic data transformations
- assign variable and value labels
- understand the difference between saving data and saving jamovi output
- copy jamovi output to a standard word processing package

## 1.7 Introduction

From the [jamovi website](#):

jamovi is a new “3rd generation” statistical spreadsheet. Designed from the ground up to be easy to use, jamovi is a compelling alternative to costly statistical products such as SPSS and SAS.

jamovi is built on top of the R statistical language, giving you access to the best the statistics community has to offer. Would you like the R code for your analyses? jamovi can provide that too.

jamovi will always be free and open - that’s one of our core values - because jamovi is made by the scientific community, for the scientific community.

The notes provided in this course will cover the basics of using jamovi: there is much more to jamovi than we will cover in this course.

## 1.8 Part 1: An introduction to jamovi


In this very brief section, we will introduce jamovi by calculating the average of six ages. Open the jamovi package in the usual way (note that while jamovi is available for Windows, MacOS and linux, most of the screenshots in these notes will be based on the macOS version.)

## 1.9 Installing jamovi

jamovi can be downloaded for no cost at <https://www.jamovi.org/download.html> for Windows, macOS and linux. I recommend downloading the “current” release (not the “solid” release). At the time of writing, Version 2.7.24 is the appropriate version to use. Download and install jamovi in the usual way.

### 1.10 A simple jamovi analysis

When you first open jamovi, it will look something like the following.



On the left-hand side of the window is the spreadsheet view, and the right is where results of statistical analyses will appear. The **spreadsheet** is where data can be entered or changed.


#### TASK

Enter the following six ages into jamovi, starting at the top-left cell, by typing each number and then hitting Enter:

20 25 23 29 21 27

If you make a mistake, simply click the incorrect cell, and enter the correct value.


Your screen should look like this:



There are two things to note here:

1. Data in jamovi are entered down a column: columns represent variables, and rows represent observations. So our six observations of age are entered in one column.
2. jamovi has given the name of A to our column of ages.

Let's rename our variable from A to Age (years). There are a number of ways of doing this: here we will click the **Data** tab (at the top of the window) which allows us to change aspects of our dataset, and then click **Setup**. By default, the first column is selected - you can choose any column simply by clicking it.




In this window, you can change many variable properties, such as variable names and variable types. To change the variable name, click the name of the variable, currently A, at the top of the window. Replace A with Age (years):

The screenshot shows the JAMOVI software interface with the 'Data' tab selected. The 'DATA VARIABLE' section is active, showing the variable 'Age (years)' with a red box around its name. The 'Measure type' is set to 'Nominal' and the 'Data type' is set to 'Integer (auto)'. The 'Missing values' field is empty. A 'Levels' list on the right shows the values 20, 21, 23, and 25. Below the 'Levels' list is a toggle for 'Retain unused levels in analyses' which is currently turned off. The data grid at the bottom shows the first six rows of data for 'Age (years)', 'B', and 'C'. The status bar at the bottom indicates 'Ready', 'Filters 0', 'Row count 6', 'Filtered 0', 'Deleted 0', 'Added 6', and 'Cells edited 18'.

|    | Age (years) | B | C |
|----|-------------|---|---|
| 1  | 20          |   |   |
| 2  | 25          |   |   |
| 3  | 23          |   |   |
| 4  | 29          |   |   |
| 5  | 21          |   |   |
| 6  | 27          |   |   |
| 7  |             |   |   |
| 8  |             |   |   |
| 9  |             |   |   |
| 10 |             |   |   |
| 11 |             |   |   |
| 12 |             |   |   |
| 13 |             |   |   |
| 14 |             |   |   |

Note that jamovi has assumed these data are *Nominal* data; this can be changed to *Continuous* by choosing the appropriate item in the **Measure type** drop-down menu:



Once you have finished naming all your variables, you can click the **Up** arrow to close the Setup tab to view the spreadsheet again.

jamovi is very flexible with its naming convention, and allows variable names that many other statistical packages would not allow. Often, the only way to get publication quality graphs or tables is to name your variables as completely as you can. However, I would recommend the following conventions when choosing variable names in jamovi:

- try to keep your variable names relatively short;
- variable names should start with a letter;
- variable names are case-sensitive (so age, Age and AGE could represent three different variables)

#### TASK

Rename the variable A with the name Age (years), and define age as a *Continuous* variable.

Now that we have entered our six ages, let's calculate the mean age. Choose **Analyses > Exploration > Descriptives**. The **Descriptives** dialog box will appear. Move the variable Age into the **Variables** box by clicking Age (years) and then clicking the right-arrow icon. The window should appear as:

The screenshot shows the JAMOVI software interface. The top menu bar includes 'Variables', 'Data', 'Analyses', and 'Edit'. Below the menu is a toolbar with icons for Exploration, T-Tests, ANOVA, Regression, Frequencies, and Factor. The main workspace is divided into two panes. The left pane, titled 'Descriptives', shows a list of variables on the left (B and C) and a 'Variables' box on the right containing 'Age (years)'. Below the 'Variables' box is a 'Split by' box. At the bottom of the left pane, there are tabs for 'Statistics' and 'Plots', and a checkbox for 'Frequency tables'. The right pane, titled 'Results', shows the 'Descriptives' table and a 'References' section.

**Descriptives**

| Age (years)        |         |
|--------------------|---------|
| N                  | 6       |
| Missing            | 0       |
| Mean               | 24.1667 |
| Median             | 24.0000 |
| Standard deviation | 3.4881  |
| Minimum            | 20      |
| Maximum            | 29      |

**References**

[1] The jamovi project (2024). *jamovi*. (Version 2.5) [Computer Software]. Retrieved from <https://www.jamovi.org>.

[2] R Core Team (2023). *R: A Language and environment for statistical computing*. (Version 4.3) [Computer software]. Retrieved from <https://cran.r-project.org>. (R packages retrieved from CRAN)

You can see the results of the analysis in the right-hand pane.

### TASK

Calculate summary statistics for Age (years) and confirm: there are 6 observations, with a mean age of 24.2 years, a standard deviation of 3.49 years, a minimum of 20 and a maximum of 29 years.

### The jamovi environment

Now that we have seen a simple example of how to use jamovi, let's describe the jamovi environment. There are two main views in jamovi: the **Spreadsheet** view, available by clicking the **Data** tab, and the **Analysis** view, available by clicking the **Analyses** tab. We will tend to use these two tabs through the course.

The unique thing about jamovi is that *all analyses are updated whenever the data are changed in the spreadsheet*. While this can be convenient, *care must be taken not to make any unintended changes to your data while in the spreadsheet view*.

jamovi also has a way of opening and saving data, using the three lines in the upper-left corner:



. This collection of commands lets you open data, save data and export data and output.

### Tip

The **Special Import** command is hardly ever used. You should use **Open** to open all types of data in jamovi.

### 1.11 Part 2: Obtaining summary statistics for continuous data

In this exercise (spanning Modules 1 and 2), we will analyse data to complete a descriptive table from a research study. The data come from a study in primary biliary cirrhosis, a condition of the liver, from Therneau and Grambsch (2010), *Modeling Survival Data: Extending the Cox Model*. By the end of this exercise, we will have completed some of Table 1.1.

Table 1.1: Summary of 418 participants from the PBC study (Therneau and Grambsch, 2000)

| Characteristic            |                      | Summary                   |
|---------------------------|----------------------|---------------------------|
| Age (years)               |                      | Mean (SD) or Median [IQR] |
| Sex                       | Male                 | n (%)                     |
|                           | Female               | n (%)                     |
| AST* (U/ml)               |                      | Mean (SD) or Median [IQR] |
| Serum bilirubin           |                      | Mean (SD) or Median [IQR] |
| Stage                     | I                    | n (%)                     |
|                           | II                   | n (%)                     |
|                           | III                  | n (%)                     |
|                           | IIIV                 | n (%)                     |
| Vital status at study end | Alive: no transplant | n (%)                     |
|                           | Alive: transplant    | n (%)                     |
|                           | Deceased             | n (%)                     |

\* aspartate aminotransferase

#### TASK

Download the table shell, saved as `PBC Table1.docx`, and the information file called `mod01_pbc_info.txt`.

### Opening a data file

Typing data directly into jamovi is not common; we usually open data that have been saved as a file. jamovi can open many types of files, including text (txt), comma separated (csv), Microsoft Excel (xlsx), R (rds), Stata (dta) and more. Here, we will open a dataset that has been stored as an R data file (which has the `.rds` suffix).

#### TASK

Load the sample data set called `mod01_pbc.rds` into jamovi using the following steps:

1. Locate the data set called `mod01_pbc.rds`. Click the file to download it, and then save it in a folder you will be able to locate later - for example, your OneDrive folder.
2. In jamovi, click the three-bar icon, then choose **Open**. jamovi usually searches in the most recently used folder, so most times you will need to **click Browse**. Browse to where you stored the dataset and click **Open**.

Confirm that there are 418 rows by examining the **Row count** at the bottom of the screen. Examine the `pbc_info.txt` file for a description of each variable.

### Assigning meaningful variable names

As we saw earlier, jamovi has can allow quite useful variable names, which will appear when creating output. For example, the variable entered as `bili` could be named `Serum bilirubin`




(mg/dl).

### TASK

Assign meaningful variable names to the variables used in Table 1. You should refer to the file `pbcc_info.txt` to determine what each variable represents.

### Summarising continuous variables

As we saw in Part 1, continuous variables can be summarised using **Analyses > Exploration**. There are three continuous variables that we would like to summarise: age, AST and serum bilirubin. Each of these can be listed in the **Exploration** dialog box, as shown below. The summaries are calculated automatically:



The screenshot shows the Jamovi software interface. The top menu bar includes 'Variables', 'Data', 'Analyses', and 'Edit'. Below the menu bar are icons for various statistical tests: Exploration, T-Tests, ANOVA, Regression, Frequencies, and Factor. The 'Analyses' menu is open, and the 'Exploration' option is selected. The 'Descriptives' dialog box is open, showing a list of variables on the left: ascites, hepato, spiders, edema, chol, albumin, copper, alkphos, and trig. The 'Variables' box on the right contains 'Age (years)', 'Serum bilirubin (mg/dL)', and 'AST (U/ml)'. The 'Split by' box is empty. The 'Descriptives' section is set to 'Variables across columns'. The 'Frequency tables' checkbox is unchecked. The 'Results' panel on the right shows the 'Descriptives' table.

|                    | Age (years) | Serum bilirubin (mg/dL) | AST (U/ml) |
|--------------------|-------------|-------------------------|------------|
| N                  | 418         | 418                     | 311        |
| Missing            | 0           | 0                       | 107        |
| Mean               | 50.7416     | 3.2208                  | 122.5068   |
| Median             | 51.0007     | 1.4000                  | 114.7000   |
| Standard deviation | 10.4472     | 4.4075                  | 56.7842    |
| Minimum            | 26.2779     | 0.3000                  | 26.3500    |
| Maximum            | 78.4394     | 28.0000                 | 457.2500   |

The 'References' section at the bottom of the Results panel contains two references:

- [1] The jamovi project (2024). *jamovi*. (Version 2.5) [Computer Software]. Retrieved from <https://www.jamovi.org>.
- [2] R Core Team (2023). *R: A Language and environment for statistical computing*. (Version 4.3) [Computer software]. Retrieved from <https://cran.r-project.org>. (R packages retrieved from CRAN snapshot 2024-01-09).

By default, the exploration command presents the number of observations, the number of missing observations, the mean, median, standard deviation, minimum and maximum. We may be interested in obtaining the interquartile range as well, so we select the **Statistics** arrow, and choose **Percentiles**:

The screenshot shows the Jamovi software interface with the 'Analyses' tab selected. In the 'Descriptives' panel, the variables 'Age (years)', 'Serum bilirubin (mg/dL)', and 'AST (U/ml)' are listed. The 'Statistics' section is expanded, showing options for Sample Size, Central Tendency, Dispersion, Mean Dispersion, Distribution, Normality, and Outliers. The 'Percentile Values' section is highlighted, showing 'Cut points for 4 equal groups' and 'Percentiles 25,50,75' selected. The 'Results' panel displays the 'Descriptives' table, which is highlighted with a red box. The table shows the following data:

|                    | Age (years) | Serum bilirubin (mg/dL) | AST (U/ml) |
|--------------------|-------------|-------------------------|------------|
| N                  | 418         | 418                     | 311        |
| Missing            | 0           | 0                       | 107        |
| Mean               | 50.7416     | 3.2208                  | 122.5068   |
| Median             | 51.0007     | 1.4000                  | 114.7000   |
| Standard deviation | 10.4472     | 4.4075                  | 56.7842    |
| Minimum            | 26.2779     | 0.3000                  | 26.3500    |
| Maximum            | 78.4394     | 28.0000                 | 457.2500   |
| 25th percentile    | 42.8323     | 0.8000                  | 80.6000    |
| 50th percentile    | 51.0007     | 1.4000                  | 114.7000   |
| 75th percentile    | 58.2409     | 3.4000                  | 151.9000   |


The 'References' panel lists two references:

- [1] The jamovi project (2024). *jamovi*. (Version 2.5) [Computer Software]. Retrieved from <https://www.jamovi.org>.
- [2] R Core Team (2023). *R: A Language and environment for statistical computing*. (Version 4.3) [Computer software]. Retrieved from <https://cran.r-project.org>. (R packages retrieved from CRAN snapshot 2024-01-09).

For each of our three continuous variables, we need to decide whether to present the mean and standard deviation, or the median and interquartile range. This decision can be made after examining a density plot (and perhaps a boxplot) for each variable.

### Producing a density plot

To produce a density plot, click the arrow next to **Plots** and choose **Density**. Plots will be produced for each variable listed in the **Variables** box. The density plots will be produced one after another, but they have been presented horizontally here:



### Producing a boxplot

Producing boxplots is done by ticking the **Box plot** box. By default, jamovi labels each of the points that it considers to be an outlier with its row number; this can be turned off if desired.

**TASK**

Obtain density plots and boxplots for age, AST and bilirubin.  
Based on these plots, decide whether the mean or the median is the appropriate summary to use for each variable.

**Saving your work from jamovi**

Now that you have made some changes to the pbc data and conducted some analyses, it is good practice to save your work. jamovi uses its own file format to save **both data and output**, using files with .omv suffix. All changes to your data will be saved, as well as all existing output. However, work saved by jamovi will only be able to be opened by jamovi - you will not easily be able to share your data or your output with colleagues who do not have jamovi. To save a jamovi session, choose **Save** in the three-lines tab.

If you want to share work with colleagues who do not have jamovi, you can use **Export** to save your data in another file format (recognising that variable and value labels will not be exported), or save your output as a pdf or html file.

**Copying output from jamovi**

An easy way to share output between your colleagues is to copy the output into a word processor package (e.g. Microsoft Word). To copy output from jamovi, you can **right-click**<sup>2</sup> the output with your mouse, and choose **Export**. This will copy the output as plain text for pasting into a Word document. If you select a single table for copying, you can also **Copy table** or **Copy table as HTML**. Whichever way you copy output into Word, you will need to make sure your output conforms with all style guides required for your final publication.

**TASK**

Complete Table 1 for continuous variables using the output generated in this exercise. You should decide on whether to present continuous variables by their means or medians, and present the most appropriate measure of spread. Include footnotes to indicate if any variables contain missing observations.

**1.12 Setting a value to missing**

As we saw in Section 1.5, it is important to explore our data to identify any unusual observations. If an error is found, the best method for correcting the error is to go back to the original data e.g. the hard copy questionnaire, to obtain the original value, entering the correct value into jamovi. If the original data is not available or the original data is also incorrect, the erroneous value is often excluded from the dataset.

Consider a sample dataset: mod01\_weight\_1000.rds, which contains the weights of 1000 people. A density plot and a boxplot should be examined before we start analysing these data:


There is a clear outlying point shown in the boxplot. Although not obvious, the same point is shown in the density plot as a small blip around 700kg. Obviously this point is unusual, and we should investigate. You will need to decide if any usual values are a data entry error or are biologically plausible. If an extreme value or “outlier”, is biologically plausible, it should be included in all analyses.

<sup>2</sup>**Windows**

- Click the right mouse button to open a context menu
- Shows options relevant to what you clicked on (copy, export, add note)

**Mac**

- Two-button mouse: Press right button
- One-button mouse/trackpad: Hold Control (Ctrl) while clicking
- Functions the same as Windows - shows context-specific options



Notice the boxed number in the boxplot: this is the record number in jamovi's spreadsheet. Click **Data** to view the spreadsheet, and scroll to record number 58:

| Variables Data Analyses |    |               |  |
|-------------------------|----|---------------|--|
| Paste<br>Clipboard      |    | Setup Compute |  |
| id                      |    | weight        |  |
| 48                      | 48 | 70.0          |  |
| 49                      | 49 | 71.7          |  |
| 50                      | 50 | 73.0          |  |
| 51                      | 51 | 62.3          |  |
| 52                      | 52 | 85.7          |  |
| 53                      | 53 | 65.2          |  |
| 54                      | 54 | 63.4          |  |
| 55                      | 55 | 72.8          |  |
| 56                      | 56 | 63.7          |  |
| 57                      | 57 | 62.4          |  |
| 58                      | 58 | 700.2         |  |
| 59                      | 59 | 72.4          |  |
| 60                      | 60 | 65.4          |  |
| 61                      | 61 | 68.1          |  |
| 62                      | 62 | 69.5          |  |

We see that there is a very high value of 700.2kg. A value as high as 700kg is likely to be a data entry error (e.g. error in entering an extra zero) and is not a plausible weight value. Here, **you should check your original data.**

If you do not have access to the original data, it would be safest to set this value as missing. You do change this in jamovi by clicking the datapoint and pressing **Delete** or **Backspace**. This has set this weight to missing.

**Note:** if an extreme value lies within the range of biological possibility it should not be set to missing.

The same process could be used to replace the incorrect value with the correct value. For example, if you do source the original medical records, you might find that the original weight was recorded in medical records as 70.2kg. We could use the same process to replace 700.2 by 70.2 by entering the correct value in the cell.

Once you have checked your data for errors, you are ready to start analysing your data.

**! Important**

Whenever you start changing data in jamovi, you should **always keep an original, unedited copy of your data**. You can do this using **Save As** to save your edited work, leaving the original data unchanged.



# An introduction to R and RStudio

## Learning outcomes

By the end of this Module, you will be able to:

- understand the difference between R and RStudio
- navigate the RStudio interface
- input and import data into R
- use R to summarise data
- perform basic data transformations
- understand the difference between saving R data and saving R output
- copy R output to a standard word processing package

## 1.13 Part 1: An introduction to R

“R is a language and environment for statistical computing and graphics.” [Link](#). It is an open-source programming language, used mainly for statistics (including biostatistics) and data science.

The aim of these notes is to introduce the R language within the RStudio environment, and to introduce the commands and procedures that are directly relevant to this course. There is so much more to R than we can cover in these notes. Relevant information will be provided throughout the course, and we will provide further references that you can explore if you are interested.

### R vs RStudio

At its heart, R is a programming language. When you install R on your computer, you are installing the language and its resources, as well as a very basic interface for using R. You can write and run R code using the basic R app, but it's not recommended.

RStudio is an “Integrated Development Environment” that runs R while also providing useful tools to help you write code and analyse data. You can think of R as an engine which does the work, and RStudio as a car that uses the engine, but also provides useful tools like GPS navigation and reversing cameras that help you drive.

**Note:** even though we recommend that you use RStudio, you still need install R. **RStudio will not run without R installed.**

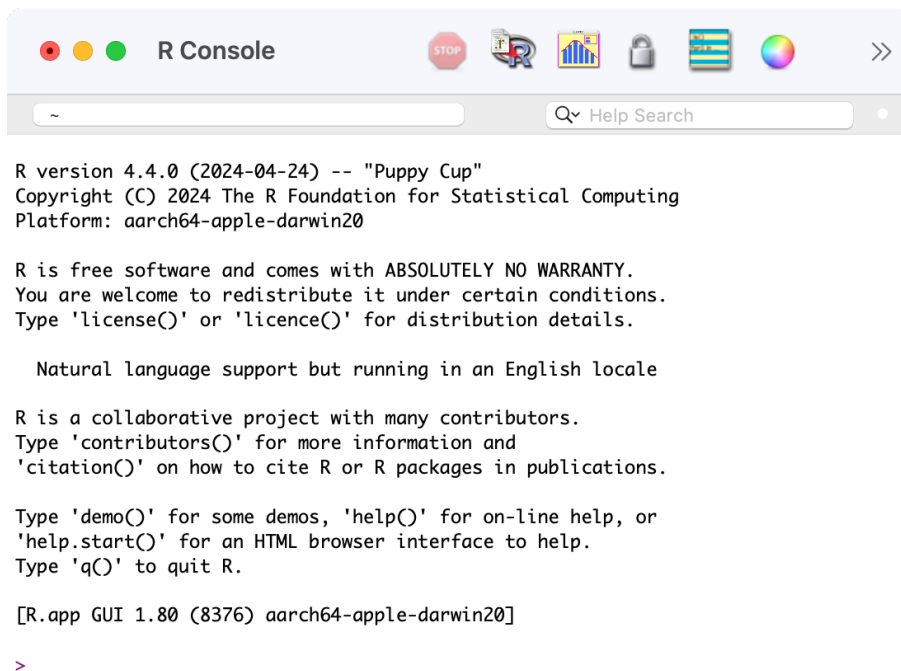
In summary, we recommend you use RStudio to write R code.

### Installing R and RStudio

#### To install R on your computer

1. Download the R installer from:
  - a. for Windows: <https://cran.r-project.org/bin/windows/base/>
  - b. for MacOS: <https://cran.r-project.org/bin/macosx/>
2. Install R by running the installer and following the installation instructions. The default settings are fine.

3. Open the R program. You should see a screen similar to below:



Near the bottom of the R screen, you will find the “>” symbol which represents R’s command line; this is where you can type your R commands for processing. If you type `1 + 2` into the command line and then hit enter you should get:

```
[1] 3
```

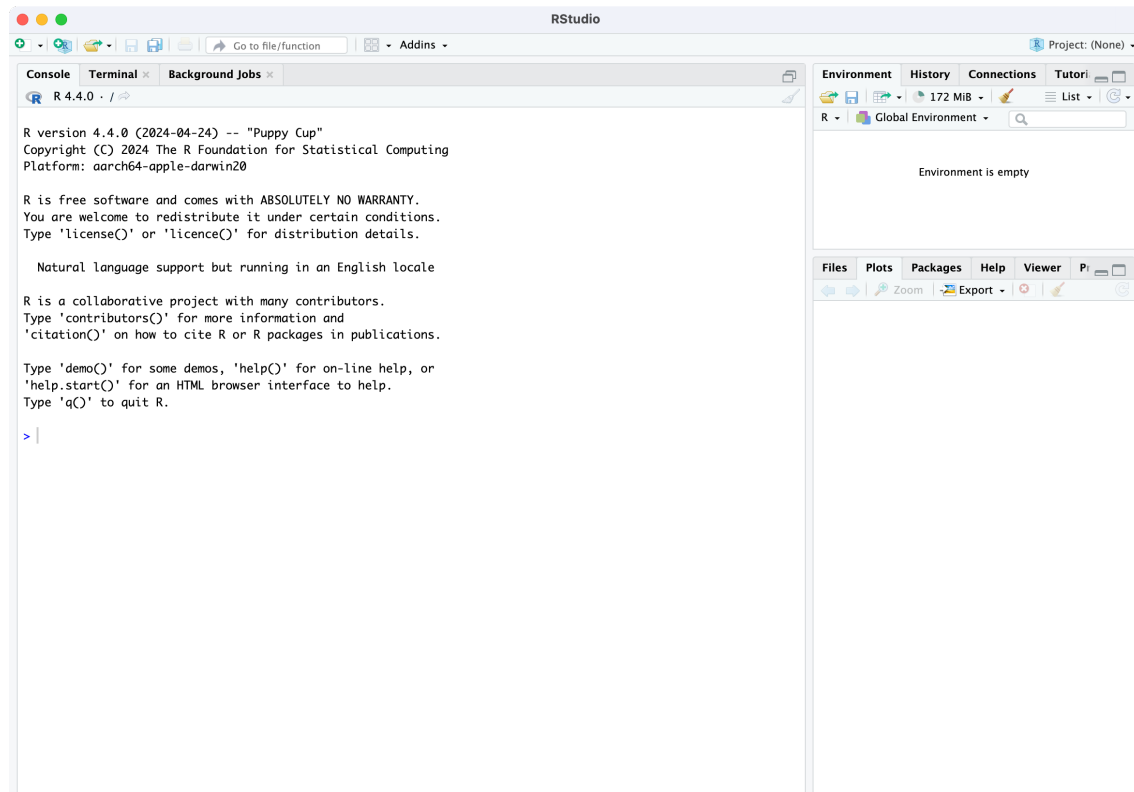
This is R performing your calculation, with the `[1]` indicating that the solution to `1 + 2` is a single number (the number 3).

At this point, close R - we will not interact with R like this in the future. You can close R by typing `quit()` at the command prompt, followed by the return key, or in the usual way of closing an application in your operating system. There is no need to save anything here if prompted.

### To install RStudio on your computer

1. Make sure you have already installed R, and verified that it is working.
2. Download the RStudio desktop installer at: <https://posit.co/download/rstudio-desktop/>. The website should detect your operating system and link to the appropriate installer for your computer.
3. Install RStudio by running the installer and following the installation instructions. The default settings are fine.
4. Open RStudio, which will appear similar to the screenshot below:





Locate the R's command line symbol ">" at the bottom of the left-hand panel. Type `1 + 2` into the command line and hit enter, and you will see:

```
[1] 3
```

This confirms that RStudio is running correctly, and can use the R language to correctly calculate the sum between 1 and 2!

RStudio currently comprises three window panes, and we will discuss these later.

#### TASK

Install R and RStudio and confirm they are both working correctly.

### Recommended setup


I will provide a recommended setup for R and RStudio in this section. You are free to use alternative workflows and setup, but this setup works well in practice.

### RStudio preferences

By default, RStudio will retain data, scripts and other objects when you quit your RStudio session. Relying on this can cause headaches, so I recommend that you set up RStudio so that it does not preserve your workspace between sessions. Open the RStudio options:

- Mac: **Edit > Settings**
- Windows: **Tools > Global Options**

Make sure you are in the **General** pane, and **untick "Restore .RData into workplace at startup"**, and choose: **"Save workspace to .RData on exit: Never"**.




## Set up a project


A project in RStudio is a folder that RStudio recognises as a place to store R scripts, data files, figures that are common to an analysis project. Setting up a folder allows much more simple navigation and specification of data files and output. More detail can be found in Chapter 8 of the excellent text: [R for Data Science](#). Using projects is not necessary, but I recommend working with projects from day one.

We will create a project called **PHCM9795** to store all the data you will use and scripts that you will write in this course. First, think about where you want to store your project folder: this could be somewhere in your *Documents* folder.


Step 1: Choose **File > New Project...** in RStudio to open the **Create Project** dialog box. Click the first option to create a project in a **New directory**:



Step 2: Click the first option: **New Project**:



Step 3: Give the project a name, by typing PHCM9795 in the "Directory name", and choose where you want to store the project by clicking the **Browse** button. Finally, click **Create** to create your project.



New Project Wizard

Back Create New Project

Directory name: 1 PHCM9795


Create project as subdirectory of: ~/Documents Browse... 2

☐ Create a git repository

☐ Use renv with this project

☐ Open in new session 3 Create Project Cancel

You will now have a new folder in your directory, which contains only one file: PHCM9795.Rproj, and the two right-hand panes of RStudio will appear as below:



#### TASK

Create a new project called PHCM9795. Store it in a folder that you will remember, so that you can store all your R work in the same folder over the term.

The top-right menu bar is showing that you are working within the PHCM9795 project, and the bottom-right window is showing the contents of that window: the single PHCM9795.Rproj file. We will add some more files to this project later.

#### A simple R analysis

In this very brief section, we will introduce R by calculating the average of six ages.

To begin, open a new R Script by choosing **File > New file > R Script**. A script (or a program) is a collection of commands that are sequentially processed by R. You can also type Ctrl+Shift+N in Windows, or Command+Shift+N in MacOS to open a new script in RStudio, or click the **New File** button at the top of the RStudio window.

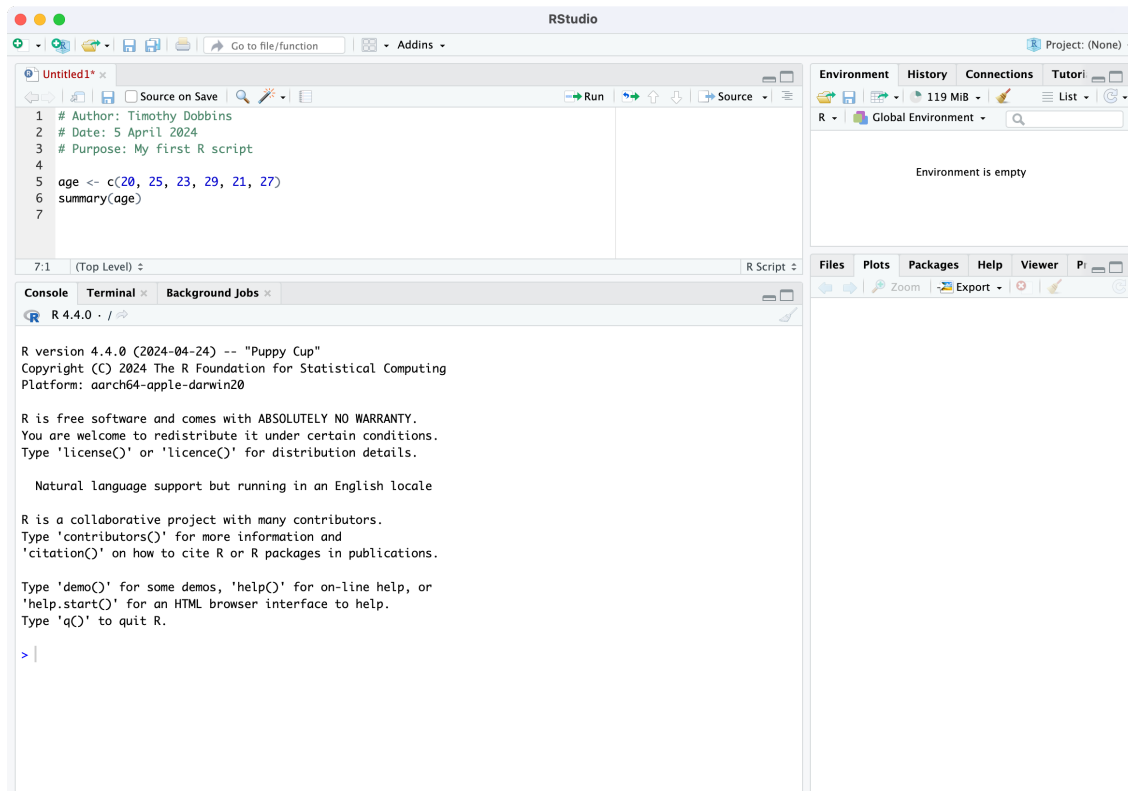
You should now see four window panes, as below. In the top-left window, type the following (replacing my name with yours, and including today's date):

```
# Author: Timothy Dobbins
# Date: 5 April 2024
# Purpose: My first R script
```

```
age <- c(20, 25, 23, 29, 21, 27)
summary(age)
```

**Note:** R is case-sensitive, so you should enter the text exactly as written in these notes.

Your screen should look like:



To run your script, choose **Code > Run Region > Run All** from the menu bar. You will see your code appear in the bottom-left window, with the following output:

```
> # Author: Timothy Dobbins
> # Date: 5 April 2024
> # Purpose: My first R script
>
> age <- c(20, 25, 23, 29, 21, 27)
>
> summary(age)
  Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
 20.00  21.50   24.00   24.17  26.50   29.00
```

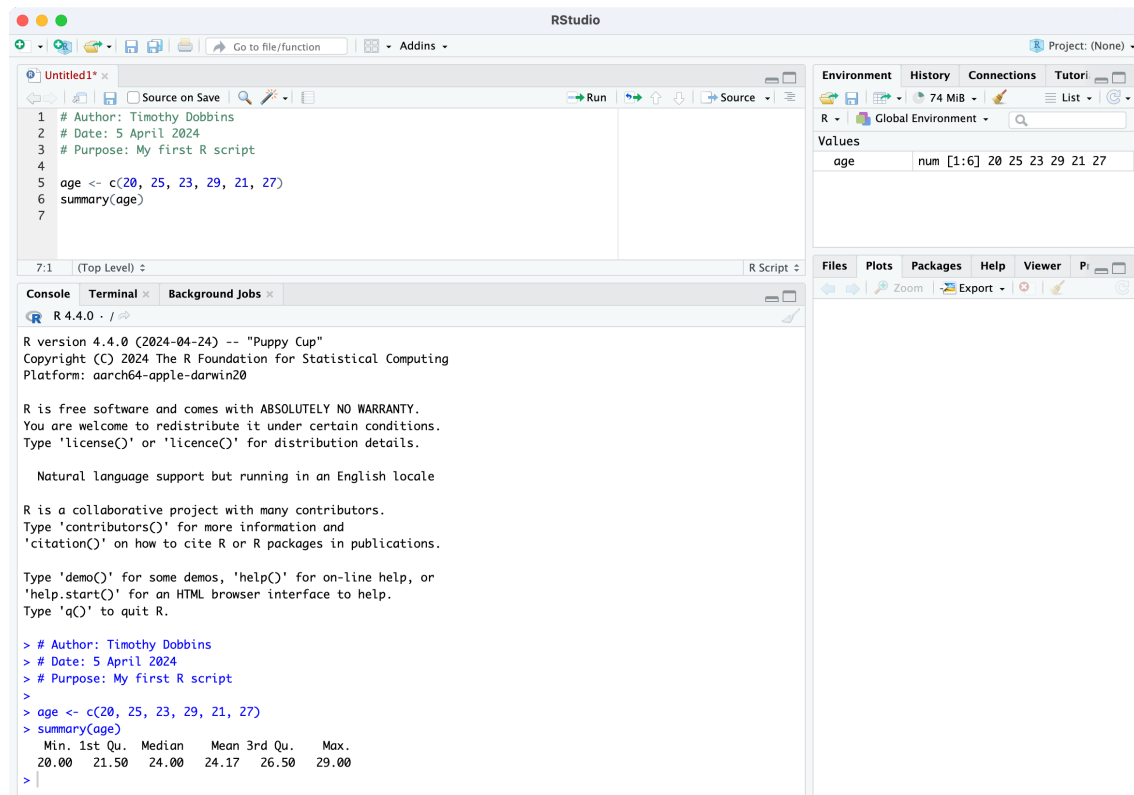
We will explain the key parts of this script later, but for now, you have entered six ages and calculated the mean age (along with five other summary statistics).

#### TASK

Type the code above into the top-left window, and run the script. Save your script within the PHCM9795 project by using **File > Save As**, using the name `my_first_analysis.R`.

## The RStudio environment

Now that we have seen a simple example of how to use R within RStudio, let's describe the RStudio environment. Let's assume that you have just run your first R script, and you have four windows as below:



The top-left window is called the **Source** window, and is where you write and edit your R scripts. Scripts can be saved by clicking **File > Save As** or by clicking on the symbol of a floppy disk at the top of the script. The file will have an extension of .R, for example script.R. Remember to give your script a meaningful title and remember to periodically save as you go.

Notice the tab at the top of the Source window which includes the name of the current file. The name of the script will be black when it has been saved, and will change to red if you have any unsaved changes.

The **Console** window, at the bottom left, contains the command line which is indicated with the symbol >. You can type commands here, but anything executed directly from the console is not saved and therefore is lost when the session ends (when you exit RStudio). You should always run your commands from a script file which you can save and use again later. When you run commands from a script, the output and any notes/errors are shown in the console. The Terminal and Jobs tabs will not be used in this course.

The **Environment** window at the top-right shows a list of objects that have been created during your session. When you close your RStudio session these objects will disappear. We will not use the History or Connections tabs in this course.

The bottom right corner contains some useful tabs, in particular the **Help** tab. When you are troubleshooting errors or learning how to use a function, the Help tab should be the first place you visit. Here you can search the help documents for all the packages you have installed. Whenever you create plots in R, these will be shown in the **Plots** tab. The **Packages** tab contains a list of installed packages and indicates which ones are currently in use (we will learn about packages later). Packages which are loaded, i.e. in use, are indicated with a tick. Some packages are in use by default when you begin a new session. You can access information about a package by clicking on its name. The **Files** tab provides a shortcut to access your files. The Viewer tab will not be used in this course.

## Some R basics

While we use R as a statistics package, R is a programming language. In order to use R effectively, we need to define some basics.

### Scripts

While R can be run completely from the command line, issuing commands one-by-one, it is most commonly run using **scripts**. A script is simply a list of commands that are processed in order. The simple analysis we conducted earlier is a very simple script. Some things to know about R scripts:

- anything appearing after a **#** is a comment, and is ignored by R. The first three lines of our script are there for ourselves (either as writers of code, or readers of code). I include comments at the beginning of each of my scripts to describe:
  - who wrote the script (useful if someone else uses your script and wants to ask questions about it);
  - when the script was written;
  - what the script does. This last point may seem odd, but it's useful to describe what this script does, and why it might differ to other scripts being used in the analysis. This is particularly useful if your scripts become long and complex.
- **R is case-sensitive**. So `age`, `AGE` and `Age` could refer to three separate variables (please don't do this!)
- use blank lines and comments to separate sections of your script

### Objects

If you do some reading about R, you may learn that R is based on **objects** and **functions**. We can think of objects as the 'nouns' of R, and functions as the 'verbs'. When we enter data, or import data into R, we are asking R to create **objects** from our data. These objects can be manipulated and transformed by **functions**, to obtain useful insights from our data.

Objects in R are created using the **assignment operator**. The most common form of the assignment operator looks like an arrow: `<-` and is typed as the `<` and `-` symbols. The simplest way of reading `<-` is as the words "is defined as". Note that it is possible to use `->` and even `=` as assignment operators, but their use is less frequent.

We created an object earlier, called `age`, which comprised six values. Let's see another example:

```
x <- 42
```

This command creates a new object called `x`, which is defined as the number 42 (or in words, "`x` is defined as 42"). Running this command gives no output in the console, but the new object appears in the top-right **Environment** panel. We can view the object in the console by typing its name:

```
# Print the object x
x
```

```
[1] 42
```

Now we see the contents of `x` in the console.

This example is rather trivial, and we rarely assign objects of just one value.



## Data structures

There are two main structures we will use to work with data in this course: **vectors** and **data frames**. A **vector** is a combination of data values, all of the same type. For example, our six ages that we entered earlier is a vector. You could think of a vector as a column of data (even though R prints vectors as rows!) And technically, even an object with only one value is a vector, a vector of size 1.

The easiest way of creating a vector in R is by using the `c()` function, where `c` stands for 'combine'. In our previous Simple Analysis in R (Section 1.13), we wrote the command:

```
age <- c(20, 25, 23, 29, 21, 27)
```

This command created a new object called `age`, and *combined* the six values of age into one vector.

Just as having a vector of size 1 is unusual, having just one column of data to analyse is also pretty unusual. The other structure we will describe here is a **data frame** which is essentially a collection of vectors, each of the same size. You could think of a data frame as being like a spreadsheet, with columns representing variables, and rows representing observations.

There are other structures in R, such as matrices and lists, which we won't discuss in this course. You may also come across the term **tibble**, which is a type of data frame.

## Functions

As the 'verbs' of R, functions transform objects. Functions can transform your data into summary statistics, graphical summaries or analysis results. For example, we used the `summary()` function to display summary statistics for our six ages.

R functions are specified by their arguments (or inputs). The arguments that can be supplied for each function can be inspected by examining the help notes for that function. To obtain help for a function, we can submit `help(summary)` (or equivalently `?summary`) in the console, or we can use the **Help** tab in the bottom-right window of RStudio. For example, the first part of the help notes for `summary` appear as:

```
summary {base}
```

R Documentation

## Object Summaries

### Description

`summary` is a generic function used to produce result summaries of the results of various model fitting functions. The function invokes particular [methods](#) which depend on the [class](#) of the first argument.

### Usage

```
summary(object, ...)
```

```
## Default S3 method:
summary(object, ..., digits, quantile.type = 7)
## S3 method for class 'data.frame'
summary(object, maxsum = 7,
        digits = max(3, getOption("digits")-3), ...)
```

```
## S3 method for class 'factor'
summary(object, maxsum = 100, ...)
```

The help notes in R can be quite cryptic, but the **Usage** section details what inputs should be specified for the function to run. Here, `summary` requires an object to be specified. In our case, we specified `age`, which is our object defined as the vector of six ages.

Most help pages also include some examples of how you might use the function. These can be found at the very bottom of the help page.

## Examples

### [Run examples](#)

```
summary(attach, digits = 4) #-> summary.data.frame(...), default precision
summary(attach $ station, maxsum = 20) #-> summary.factor(...)

lst <- unclass(attach$station) > 20 # logical with NAs
## summary.default() for logicals -- different from *.factor:
summary(lst)
summary(as.factor(lst))
```

The `summary()` function is quite simple, in that it only requires one input, the object to be summarised. More complex functions might require a number of inputs. For example, the help notes for the `descriptives()` function in the `jmv` package show a large number of inputs can be specified. Instructions for installing the `jmv` package will be provided later, this help-screen is included for illustration only.

`descriptives(jmv)`

R Documentation

## Descriptives

### Description

Descriptives are an assortment of summarising statistics, and visualizations which allow exploring the shape and distribution of data. It is good practice to explore your data with descriptives before proceeding to more formal tests.

### Usage

```
descriptives(data, vars, splitBy = NULL, freq = FALSE,
  desc = "columns", hist = FALSE, dens = FALSE, bar = FALSE,
  barCounts = FALSE, box = FALSE, violin = FALSE, dot = FALSE,
  dotType = "jitter", boxMean = FALSE, boxLabelOutliers = TRUE,
  qq = FALSE, n = TRUE, missing = TRUE, mean = TRUE,
  median = TRUE, mode = FALSE, sum = FALSE, sd = TRUE,
  variance = FALSE, range = FALSE, min = TRUE, max = TRUE,
  se = FALSE, ci = FALSE, ciWidth = 95, iqr = FALSE,
  skew = FALSE, kurt = FALSE, sw = FALSE, pcEqGr = FALSE,
  pcNEqGr = 4, pc = FALSE, pcValues = "25,50,75", formula)
```

There are two things to note here. First, notice that the first two inputs are listed with `no =` symbol, but all other inputs are listed with `=` symbols (with values provided after the `=` symbol). This means that everything apart from `data` and `vars` have **default** values. We are free to not specify values for these inputs if we are happy with the defaults provided. For example, by default the variance is not calculated (as `variance = FALSE`). To obtain the variance as well as the standard deviation, we can change this default to `variance = TRUE`:

```
# Only the standard deviation is provided as the measure of variability
descriptives(data = pbc, vars = age)
```

```
# Additionally request the variance to be calculated
descriptives(data = pbc, vars = age, variance = TRUE)
```

Second, for functions with multiple inputs, we can specify the input name and its value, or we can ignore the input name and specify just the input values **in the order listed in the Usage section**. So the following are equivalent:

```
# We can specify that the dataset to be summarised is pbc,
# and the variable to summarise is age:
descriptives(data = pbc, vars = age)
```

```
# We can omit the input name, as long as we keep the inputs in the correct order -
# that is, dataset first, variable second:
descriptives(pbc, age)
```

```
# We can change the order of the inputs, as long as we specify the input name:
descriptives(vars = age, data = pbc)
```

In this course, we will usually provide all the input names, even when they are not required. As you become more familiar with R, you will start to use the shortcut method.

### The curse of inconsistency

As R is an open-source project, many people have contributed to its development. This has led to a frustrating part of R: some functions require a single object to be specified, but some require you to specify a data frame and select variables for analysis. Let's see an example.

The help for `summary()` specifies the usage as: `summary(object, ...)`. This means we need to specify a single object to be summarised. An object could be a single column of data (i.e. a vector), or it could be a data frame. If we have a data frame called `pbc` which contains many variables, the command `summary(pbc)` would summarise every variable in the data frame.

What if we only wanted to summarise the age of the participants in the data frame? To select a single variable from a data frame, we can use the following syntax: `dataframe$variable`. So to summarise just age from this data frame, we would use: `summary(pbc$age)`.

Compare this with the `descriptives()` function in the `jmv` package. We saw earlier that the two required inputs for `descriptives()` are `data` (the data frame to be analysed) and `vars` (the variables to be analysed). So to summarise age from the `pbc` data frame, we would specify `descriptives(data=pbc, vars=age)`.

This inconsistency will seem maddening at first, and will continue to be maddening! Reading the **usage** section of the help pages is a useful way to determine whether you should specify an object (like `pbc$age`) or a data frame and a list of variables.

### Packages

A **package** is a collection of functions, documentation (and sometimes datasets) that extend the capabilities of R. Packages have been written by R users to be freely distributed and used by others. R packages can be obtained from many sources, but the most common source is CRAN: the Comprehensive R Archive Network.

A useful way of thinking about R is that R is like a smartphone, with packages being like apps which are downloaded from CRAN (similar to an app-store). When you first install R, it comes with a basic set of packages (apps) installed. You can do a lot of things with these basic packages, but sometimes you might want to do things differently, or you may want to perform some analyses that can't be done using the default packages. In these cases, you can install a package.

Like installing an app on a smartphone, you only need to *install* a package once. But each time you want to use the package, you need to *load* the package into R.


### How to install a package

There are a couple of ways to install a package. You can use the `install.packages()` function if you know the exact name of the package. Let's use an example of installing the `skimr` package, which gives a very nice, high-level overview of any data frame. We can install `skimr` by typing the following into the console:

```
install.packages("skimr")
```

Note the use of the quotation marks.

Alternatively, RStudio offers a graphical way of installing packages that can be accessed via **Tools > Install Packages**, or via the **Install** button at the top of the **Packages** tab in the bottom-right window. You can begin typing the name of the package in the dialog box that appears, and RStudio will use predictive text to offer possible packages:



While writing code is usually the recommended way to use R, installing packages is an exception. Using **Tools > Install Packages** is perfectly fine, because you only need to install a package once.

### How to load a package

When you begin a new session in RStudio, i.e. when you open RStudio, only certain core packages are automatically loaded. You can use the `library()` function to load a package that has previously been installed. For example, now that we have installed `skimr`, we need to load it before we can use it:

```
library(skimr)
```

Note that quotation marks are not required for the `library()` function (although they can be included if you really like quotation marks!). Another example of R's inconsistency!!!

#### TASK

Install the packages `jmv` and `skimr` using **Tools > Install packages**, or by typing into the console:

```
install.packages("jmv")
install.packages("skimr")
```

### Installing vs loading packages

Package installation:

- use the `install.packages()` function (note the 's') or **Tools > Install packages**
- the package name must be surrounded by quotation marks
- only needs to be done once

Package loading:

- use the `library()` function
- the package name does not need to be surrounded by quotation marks
- must be done for each R session

### What is this thing called the tidyverse?

If you have done much reading about R, you may have come across the tidyverse:

“The tidyverse is an opinionated collection of R packages designed for data science. All packages share an underlying design philosophy, grammar, and data structures.”  
<https://www.tidyverse.org/>

Packages in the tidyverse have been designed with a goal to make using R more consistent by defining a “grammar” to manipulate data, examine data and draw conclusions from data. The tidyverse is a powerful set of packages, and we will use only a small part of its capabilities.

More information about the tidyverse can be found using freely available resources such as: [Tidyverse Skills for Data Science](#) and [R for Data Science](#).

### 1.14 Part 2: Obtaining summary statistics for continuous data

In this exercise (spanning Modules 1 and 2), we will analyse data to complete a descriptive table from a research study. The data come from a study in primary biliary cirrhosis, a condition of the liver, from Modeling Survival Data: Extending the Cox Model by Therneau and Grambsch (2010). By the end of this exercise, we will have completed some of Table 1.2.

Table 1.2: Summary of 418 participants from the PBC study (Therneau and Grambsch, 2000)

| Characteristic            |                      | Summary                   |
|---------------------------|----------------------|---------------------------|
| Age (years)               |                      | Mean (SD) or Median [IQR] |
| Sex                       | Male                 | n (%)                     |
|                           | Female               | n (%)                     |
| AST* (U/ml)               |                      | Mean (SD) or Median [IQR] |
| Serum bilirubin           |                      | Mean (SD) or Median [IQR] |
| Stage                     | I                    | n (%)                     |
|                           | II                   | n (%)                     |
|                           | III                  | n (%)                     |
|                           | IIIV                 | n (%)                     |
| Vital status at study end | Alive: no transplant | n (%)                     |
|                           | Alive: transplant    | n (%)                     |
|                           | Deceased             | n (%)                     |

\* aspartate aminotransferase

This table is available in `Table1.docx`.


#### TASK

Download the table shell, saved as `PBC Table1.docx`, and the information file called `mod01_pbc_info.txt`.

### Set up your data

We created a project in the previous step. We will now create a folder to store all the data for this course. Storing the data within the project makes life much easier!

Create a new folder by clicking the **New Folder** icon in the **Files** tab at the bottom-right:




Call the new folder **data**.

Click on this folder to open it, and then create two new folders: **activities** and **examples**.

Download the “Data sets: for learning activities”, and use Windows Explorer or MacOS Finder to save these data sets in **activities**. Save the “Data sets: example data from course notes” into the **examples** folder.

Your **activities** folder should look like:



Click the two dots next to the up-arrow at the top of the folder contents to move back up the folder structure. Note that you need to click the dots, and not the up-facing green arrow!

## Reading a data file

Typing data directly into R is not common; we usually read data that have been previously saved. In this example, we will read an `.rds` file using the `readRDS()` function, which has only one input: the location of the file.

### TASK

- 1 - Confirm that the `mod01_pdc.rds` file is in the `activities` sub-folder within the `data` folder (as per the previous steps).
- 2 - Load the `skimr` package, and use the `readRDS()` function to read the file into R, assigning it to a data frame called `pbc`. Because we set up our project, we can locate our data easily

by telling R to use the file: "data/activities/mod01\_pdc.rds", which translates as: the file mod01\_pdc.rds which is located in the activities sub-folder within the data folder.

```
library(skimr)
```

```
pbc <- readRDS("data/activities/mod01_pbc.rds")
```

3 - We can now use the `summary()` function to examine the pbc dataset:

```
summary(pbc)
```

| id              | time            | status           | trt             |
|-----------------|-----------------|------------------|-----------------|
| Min. : 1.0      | Min. : 41       | Min. : 0.0000    | Min. : 1.000    |
| 1st Qu.: 105.2  | 1st Qu.: 1093   | 1st Qu.: 0.0000  | 1st Qu.: 1.000  |
| Median : 209.5  | Median : 1730   | Median : 0.0000  | Median : 1.000  |
| Mean : 209.5    | Mean : 1918     | Mean : 0.8301    | Mean : 1.494    |
| 3rd Qu.: 313.8  | 3rd Qu.: 2614   | 3rd Qu.: 2.0000  | 3rd Qu.: 2.000  |
| Max. : 418.0    | Max. : 4795     | Max. : 2.0000    | Max. : 2.000    |
|                 |                 |                  | NAs : 106       |
| age             | sex             | ascites          | hepato          |
| Min. : 26.28    | Min. : 1.000    | Min. : 0.00000   | Min. : 0.0000   |
| 1st Qu.: 42.83  | 1st Qu.: 2.000  | 1st Qu.: 0.00000 | 1st Qu.: 0.0000 |
| Median : 51.00  | Median : 2.000  | Median : 0.00000 | Median : 1.0000 |
| Mean : 50.74    | Mean : 1.895    | Mean : 0.07692   | Mean : 0.5128   |
| 3rd Qu.: 58.24  | 3rd Qu.: 2.000  | 3rd Qu.: 0.00000 | 3rd Qu.: 1.0000 |
| Max. : 78.44    | Max. : 2.000    | Max. : 1.00000   | Max. : 1.0000   |
|                 |                 | NAs : 106        | NAs : 106       |
| spiders         | edema           | bili             | chol            |
| Min. : 0.0000   | Min. : 0.0000   | Min. : 0.300     | Min. : 120.0    |
| 1st Qu.: 0.0000 | 1st Qu.: 0.0000 | 1st Qu.: 0.800   | 1st Qu.: 249.5  |
| Median : 0.0000 | Median : 0.0000 | Median : 1.400   | Median : 309.5  |
| Mean : 0.2885   | Mean : 0.1005   | Mean : 3.221     | Mean : 369.5    |
| 3rd Qu.: 1.0000 | 3rd Qu.: 0.0000 | 3rd Qu.: 3.400   | 3rd Qu.: 400.0  |
| Max. : 1.0000   | Max. : 1.0000   | Max. : 28.000    | Max. : 1775.0   |
| NAs : 106       |                 |                  | NAs : 134       |
| albumin         | copper          | alkphos          | ast             |
| Min. : 1.960    | Min. : 4.00     | Min. : 289.0     | Min. : 26.35    |
| 1st Qu.: 3.243  | 1st Qu.: 41.25  | 1st Qu.: 871.5   | 1st Qu.: 80.60  |
| Median : 3.530  | Median : 73.00  | Median : 1259.0  | Median : 114.70 |
| Mean : 3.497    | Mean : 97.65    | Mean : 1982.7    | Mean : 122.56   |
| 3rd Qu.: 3.770  | 3rd Qu.: 123.00 | 3rd Qu.: 1980.0  | 3rd Qu.: 151.90 |
| Max. : 4.640    | Max. : 588.00   | Max. : 13862.4   | Max. : 457.25   |
|                 | NAs : 108       | NAs : 106        | NAs : 106       |
| trig            | platelet        | protime          | stage           |
| Min. : 33.00    | Min. : 62.0     | Min. : 9.00      | Min. : 1.000    |
| 1st Qu.: 84.25  | 1st Qu.: 188.5  | 1st Qu.: 10.00   | 1st Qu.: 2.000  |
| Median : 108.00 | Median : 251.0  | Median : 10.60   | Median : 3.000  |
| Mean : 124.70   | Mean : 257.0    | Mean : 10.73     | Mean : 3.024    |
| 3rd Qu.: 151.00 | 3rd Qu.: 318.0  | 3rd Qu.: 11.10   | 3rd Qu.: 4.000  |
| Max. : 598.00   | Max. : 721.0    | Max. : 18.00     | Max. : 4.000    |
| NAs : 136       | NAs : 11        | NAs : 2          | NAs : 6         |

An alternative to the `summary()` function is the `skim()` function in the `skimr` package, which produces summary statistics as well as rudimentary histograms:

```
skim(pbc)
```

| Data Summary           |           |               |        |       |      |       |      |       |        |      |
|------------------------|-----------|---------------|--------|-------|------|-------|------|-------|--------|------|
| Name                   | pbc       |               |        |       |      |       |      |       |        |      |
| Number of rows         | 418       |               |        |       |      |       |      |       |        |      |
| Number of columns      | 20        |               |        |       |      |       |      |       |        |      |
| Column type frequency: |           |               |        |       |      |       |      |       |        |      |
| numeric                | 20        |               |        |       |      |       |      |       |        |      |
| Group variables        |           |               |        |       |      |       |      |       |        |      |
| None                   |           |               |        |       |      |       |      |       |        |      |
| Variable type: numeric |           |               |        |       |      |       |      |       |        |      |
| skim_variable          | n_missing | complete_rate | mean   | sd    | p0   | p25   | p50  | p75   | p100   | hist |
| 1 id                   | 0         | 1             | 210.   | 121.  | 1    | 105.  | 210. | 314.  | 418    |      |
| 2 time                 | 0         | 1             | 1918.  | 1105. | 41   | 1093. | 1730 | 2614. | 4795   |      |
| 3 status               | 0         | 1             | 0.830  | 0.956 | 0    | 0     | 0    | 2     | 2      |      |
| 4 trt                  | 106       | 0.746         | 1.49   | 0.501 | 1    | 1     | 1    | 2     | 2      |      |
| 5 age                  | 0         | 1             | 50.7   | 10.4  | 26.3 | 42.8  | 51.0 | 58.2  | 78.4   |      |
| 6 sex                  | 0         | 1             | 1.89   | 0.307 | 1    | 2     | 2    | 2     | 2      |      |
| 7 ascites              | 106       | 0.746         | 0.0769 | 0.267 | 0    | 0     | 0    | 0     | 1      |      |
| 8 hepato               | 106       | 0.746         | 0.513  | 0.501 | 0    | 0     | 1    | 1     | 1      |      |
| 9 spiders              | 106       | 0.746         | 0.288  | 0.454 | 0    | 0     | 0    | 1     | 1      |      |
| 10 edema               | 0         | 1             | 0.100  | 0.253 | 0    | 0     | 0    | 0     | 1      |      |
| 11 bili                | 0         | 1             | 3.22   | 4.41  | 0.3  | 0.8   | 1.4  | 3.4   | 28     |      |
| 12 chol                | 134       | 0.679         | 370.   | 232.  | 120  | 250.  | 310. | 400   | 1775   |      |
| 13 albumin             | 0         | 1             | 3.50   | 0.425 | 1.96 | 3.24  | 3.53 | 3.77  | 4.64   |      |
| 14 copper              | 108       | 0.742         | 97.6   | 85.6  | 4    | 41.2  | 73   | 123   | 588    |      |
| 15 alkphos             | 106       | 0.746         | 1983.  | 2140. | 289  | 872.  | 1259 | 1980  | 13862. |      |
| 16 ast                 | 106       | 0.746         | 123.   | 56.7  | 26.4 | 80.6  | 115. | 152.  | 457.   |      |
| 17 trig                | 136       | 0.675         | 125.   | 65.1  | 33   | 84.2  | 108  | 151   | 598    |      |
| 18 platelet            | 11        | 0.974         | 257.   | 98.3  | 62   | 188.  | 251  | 318   | 721    |      |
| 19 protime             | 2         | 0.995         | 10.7   | 1.02  | 9    | 10    | 10.6 | 11.1  | 18     |      |
| 20 stage               | 6         | 0.986         | 3.02   | 0.882 | 1    | 2     | 3    | 4     | 4      |      |

The `summary()` and `skim()` functions are useful to give a quick overview of a dataset: how many variables are included, how variables are coded, which variables contain missing data and a crude histogram showing the distribution of numeric variables.

### Summarising continuous variables

One of the most flexible functions for summarising continuous variables is the `descriptives()` function from the `jmv` package. The function is specified as `descriptives(data=, vars=)` where:

- `data` specifies the dataframe to be analysed
- `vars` specifies the variable(s) of interest, with multiple variables combined using the `c()` function

We can summarise the three continuous variables in the `pbc` data: `age`, `AST` and serum bilirubin, as shown below.

```
library(jmv)
```

```
descriptives(data = pbc, vars = c(age, ast, bili))
```

DESCRIPTIVES

Descriptives

|                    | age      | ast      | bili      |
|--------------------|----------|----------|-----------|
| N                  | 418      | 312      | 418       |
| Missing            | 0        | 106      | 0         |
| Mean               | 50.74155 | 122.5563 | 3.220813  |
| Median             | 51.00068 | 114.7000 | 1.400000  |
| Standard deviation | 10.44721 | 56.69952 | 4.407506  |
| Minimum            | 26.27789 | 26.35000 | 0.3000000 |



|         |          |          |          |
|---------|----------|----------|----------|
| Maximum | 78.43943 | 457.2500 | 28.00000 |
|---------|----------|----------|----------|

By default, the `descriptives` function presents the mean, median, standard deviation, minimum and maximum. We can request additional statistics, such as the quartiles (which are called the percentiles, or `pc`, in the `descriptives` function):

```
descriptives(data = pbc, vars = c(age, ast, bili), pc = TRUE)
```

#### DESCRIPTIVES

##### Descriptives

|                    | age      | ast      | bili      |
|--------------------|----------|----------|-----------|
| N                  | 418      | 312      | 418       |
| Missing            | 0        | 106      | 0         |
| Mean               | 50.74155 | 122.5563 | 3.220813  |
| Median             | 51.00068 | 114.7000 | 1.400000  |
| Standard deviation | 10.44721 | 56.69952 | 4.407506  |
| Minimum            | 26.27789 | 26.35000 | 0.3000000 |
| Maximum            | 78.43943 | 457.2500 | 28.00000  |
| 25th percentile    | 42.83231 | 80.60000 | 0.8000000 |
| 50th percentile    | 51.00068 | 114.7000 | 1.400000  |
| 75th percentile    | 58.24093 | 151.9000 | 3.400000  |

#### Producing a density plot



We can add `dens=TRUE` to the `descriptives` function to produce a density plot for each listed variable:


```
descriptives(data = pbc, vars = c(age, ast, bili), pc = TRUE, dens = TRUE)
```

#### DESCRIPTIVES


##### Descriptives

|                    | age      | ast      | bili      |
|--------------------|----------|----------|-----------|
| N                  | 418      | 312      | 418       |
| Missing            | 0        | 106      | 0         |
| Mean               | 50.74155 | 122.5563 | 3.220813  |
| Median             | 51.00068 | 114.7000 | 1.400000  |
| Standard deviation | 10.44721 | 56.69952 | 4.407506  |
| Minimum            | 26.27789 | 26.35000 | 0.3000000 |
| Maximum            | 78.43943 | 457.2500 | 28.00000  |
| 25th percentile    | 42.83231 | 80.60000 | 0.8000000 |
| 50th percentile    | 51.00068 | 114.7000 | 1.400000  |
| 75th percentile    | 58.24093 | 151.9000 | 3.400000  |






Note that the density plots are plotted separately in the **Plot** window. They can be viewed using the arrows at the top of the **Plot** window:



A more flexible way of constructing a density plot is by using the `plot()` function within R, using the syntax: `plot(density(dataframe$variable))`, which plots the variable from the dataframe. For example, the default density plot for the age column of the pbc data:


```
plot(density(pbc$age))
```



This plot can be improved by using `xlab=" "` and `main=" "` to assign labels for the x-axis and overall title respectively:

```
plot(
  density(pbc$age),
  xlab = "Age in years",
  main = "Density plot of participant age from pbc study data"
)
```

### Density plot of participant age from pbc study data



**Note**

The `density()` function requires the analysis variable to contain no missing values, and will give an error if there are any missing values. We can use the option `na.rm=TRUE` to request that the density function ignore any missing values. For example:


```
plot(density(pbc$ast, na.rm=TRUE),
     xlab="AST (U/mL)",
     main="Density plot of aspartate aminotransferase from pbc study data")
```

**Producing a boxplot**

Like the density plot, boxplots can be requested in the `descriptives` function by using `box=TRUE`.

The `boxplot` function is an alternative, more flexible function, again specifying the dataframe to use and the variable to be plotted as `dataframe$variable`. Labels can be applied in the same way as the histogram:

```
boxplot(
  pbc$age,
  xlab = "Age (years)",
  main = "Boxplot of participant age from pbc study data"
)
```

**Boxplot of participant age from pbc study data****TASK**

Obtain density plots and boxplots for age, AST and bilirubin. Based on these plots, decide whether the mean or the median is the appropriate summary to use for each variable.

**Saving data in R**

There are many ways to save data from R, depending on the type of file you want to save. The recommendation for this course is to save your data using the `.rds` format, using the `saveRDS()` function, which takes two inputs: `saveRDS(object, file)`. Here, `object` is the R object to be

saved (usually a data frame), and `file` is the location for the file to be saved (file name and path, including the `.rds` suffix).

It is not necessary to save our PBC data, as we have made only minor changes to the data that can be replicated by rerunning our script. If you had made major changes and wanted to save your data, you could use:

```
saveRDS(pbc, file="pbc_revised.rds")
```

### Copying output from R

It is important to note that saving your data or your script in R will not save your output. The easiest way to retain the output of your analyses is to copy the output from the Console into a word processor package (e.g. Microsoft Word) before closing R.

Unfortunately, by default, R is not ideal for creating publication quality tables. There are many packages that will help in this process, such as R Markdown, `huxtable`, `gt` and `gtsummary`, but their use is beyond the scope of this course. [R Markdown for Scientists](#) provides an excellent introduction to R Markdown.

#### TASK


Complete Table 1 for continuous variables using the output generated in this exercise. You should decide on whether to present continuous variables by their means or medians, and present the most appropriate measure of spread. Include footnotes to indicate if any variables contain missing observations.

### 1.15 Setting a value to missing


As we saw in Section 1.5, it is important to explore our data to identify any unusual observations. If an error is found, the best method for correcting the error is to go back to the original data e.g. the hard copy questionnaire, to obtain the original value, entering the correct value into R. If the original data is not available or the original data is also incorrect, the erroneous value is often excluded from the dataset.

Consider a sample dataset: `mod01_weight_1000.rds`, which contains the weights of 1000 people. A density plot and a boxplot should be examined before we start analysing these data:

**Density plot of 1000 weight**



**Boxplot of 1000 weights**



There is a clear outlying point shown in the boxplot. Although not obvious, the same point is shown in the density plot as a small blip around 700kg. Obviously this point is unusual, and we should investigate.

We can view any outlying observations in the dataset using the `subset` function. You will need to decide if these values are a data entry error or are biologically plausible. If an extreme value or “outlier”, is biologically plausible, it should be included in all analyses.

For example, to list any observations from the `sample` dataset with a weight larger than 200:

```
subset(sample, weight > 200)
```

```
   id weight
58 58  700.2
```

We see that there is a very high value of 700.2kg. A value as high as 700kg is likely to be a data entry error (e.g. error in entering an extra zero) and is not a plausible weight value. Here, **you should check your original data**.

If you do not have access to the original data, it would be safest to set this value as missing. You do change this in R by using an `ifelse` statement to recode the incorrect weight of 700.2kg into a missing value. **A missing value in R is represented by NA.**

The form of the `ifelse` statement is as follows: `ifelse(test, value_if_true, value_if_false)`

We will write code to:

- create a new column (called `weight_clean`) in the `sample` dataframe (i.e. `sample$weight_clean`)
- test whether `weight` is equal to 700.2
  - if this is true, we will assign `weight_clean` to be NA
  - otherwise `weight_clean` will equal the value of `weight`

Putting it all together:

```
sample$weight_clean = ifelse(sample$weight == 700.2, NA, sample$weight)
```

**Note:** if an extreme value lies within the range of biological possibility it should not be set to missing.

The same syntax could be used to replace the incorrect value with the correct value. For example, if you do source the original medical records, you might find that the original weight was recorded in medical records as 70.2kg. We could use similar syntax to replace 700.2 by 70.2:

```
sample$weight_clean = ifelse(sample$weight==700.2, 70.2, sample$weight)
```

Once you have checked your data for errors, you are ready to start analysing your data.

**What on earth: == ?**

In R, the test of equality is denoted by two equal signs: `==`. So we would use `==` to test whether an observation is equal to a certain value. Let's see an example:

```
# Test whether 6 is equal to 6
6 == 6
```

```
[1] TRUE
```

```
# Test whether 6 is equal to 42  
6 == 42
```

```
[1] FALSE
```

You can read the `==` as “is equal to”. So the code `sample$weight == 700.2` is read as: “is the value of weight from the data frame sample equal to 700.2?”. In our `ifelse` statement above, if this condition is true, we replace `weight` by 70.2; if it is false, we leave `weight` as is.



# Activities

## Activity 1.1

25 participants were enrolled in a 3-week weight loss program. Table 1.3 presents the weight loss (in grams) of the participants.

Table 1.3: Weight loss (g) for 25 participants

|     |     |     |     |     |
|-----|-----|-----|-----|-----|
| 255 | 198 | 283 | 312 | 283 |
| 57  | 85  | 312 | 142 | 113 |
| 227 | 283 | 255 | 340 | 142 |
| 113 | 312 | 227 | 85  | 170 |
| 255 | 198 | 113 | 227 | 255 |

- These data have been saved as `Activity_1.1.rds`. Read the data into your software package.
- What type of data are these?
- Construct an appropriate graph to display the distribution of participants' weight loss. Provide appropriate labels for the axes and give the graph an appropriate title.

## Activity 1.2

Which of the following statements are true? The more dispersed, or spread out, a set of observations are:

- The smaller the mean value
- The larger the standard deviation
- The smaller the variance

## Activity 1.3

A study was conducted to investigate the general health of undergraduate students at Australian universities. A random selection of 218 UNSW undergraduate students from UNSW were enrolled, and their diastolic blood pressures are provided in the dataset `Activity_1.3.csv`. Compute the mean, median, range and standard deviation of diastolic blood pressure.

## Activity 1.4

The ages of 100 study participants have been saved as `Activity_1.4.rds`. Estimate the:

- mean and median;
- standard deviation and interquartile range;
- range.

Plot the data. Is there anything unusual about the ages? What do you think is a possible explanation for this?


A clean version of the data have been saved as `Activity_1.4_clean.rds`. Recalculate the summary statistics and recreate the plots using the clean data.

Based on this exercise, what is your advice on coding unusual or missing values in data?

# Appendix



# Analysis flowchart





# References

- Australian Bureau of Statistics. 2024. *Causes of Death, Australia, 2023*.  
<https://www.abs.gov.au/statistics/health/causes-death/causes-death-australia/latest-release>.
- Australian Institute of Health and Welfare. 2024. "Australia's Mothers and Babies." In *Australian Institute of Health and Welfare*. <https://www.aihw.gov.au/reports/mothers-babies/australias-mothers-babies/contents/about>.
- Australian Institute of Health and Welfare. 2025. *Australia's Health*.  
<https://www.aihw.gov.au/reports-data/australias-health>.
- Bland, Martin. 2015. *An Introduction to Medical Statistics*. 4th Edition. Oxford University Press.
- Kirkwood, Betty, and Jonathan Sterne. 2001. *Essentials of Medical Statistics*. 2nd edition. Wiley-Blackwell.
- Therneau, Terry M., and Patricia M. Grambsch. 2010. *Modeling Survival Data: Extending the Cox Model*. Springer.

